

The title is framed by decorative white dashed lines and arrows on a blue grid background. A horizontal dashed line with arrowheads at both ends is positioned above the text. A vertical dashed line with arrowheads at both ends is positioned to the right of the text. A curved dashed arrow in the top right corner points from the horizontal line towards the vertical line. A curved dashed arrow in the bottom left corner points from the vertical line towards the horizontal line.

Technical dive into Ethereum

Successful miner's computer

Takes the following steps and broadcasts block header, H_t , to network

Determine Transactions

Miner picks transactions to process (from those broadcasted)

Determine Ommers

Miner finds and includes valid ommers

Apply Rewards

Update account balance(s) to reward valid blocks

Compute a Valid State

Block Finalisation Defines result of all selected state transitions

State Transition Cycle Defines result of a single transaction $\sigma_{t+1} = Y(\sigma_t, T)$

Execution Cycle Defines result of a single cycle of the state machine

Ethereum Virtual Machine, EVM

Instruction Set

0x00 STOP
0x01 ADD
0x02 MUL
0x03 SUB
0x04 DIV
...

Execution Environment, I

Code owner, I_c

sender, I_s

Gas price, I_g

Input data, I_d

causer, I_c

value, I_v

Machine code, I_m

Block header, H_t

Message-call depth, I_d

Substate, A

Suicide Set, A_s

Log Series, A_l

Refund Balance, A_r

World State, σ

Addresses

Account States

Balances

Storage and Code

Machine state, μ

Gas available, g

Program counter, pc

Memory contents, m

No words, n

Stack contents, s

Iterator Function, O

Get next instruction from I_c

Get items to add/remove from stack, $\Delta = \alpha \div \beta$

Update machine stack

Subtract gas used

Increment Program counter

Halt?


Yes

No

All transactions?

Yes

Proof of Work



Mining Network
"Oh look, some transactions have been broadcast to the network. Let's race each other to create a new valid block... GO"

Ommers (or Uncles)
Valid, yet redundant blocks where $PoW(H_{om}, n_{om}, d) < 2^{256}/H_t$

Ethereum Blockchain Mechanism (Proof Of Work)

An interpretation of the Ethereum Project Yellow Paper

G. Wood, "Ethereum: A secure decentralised generalised transaction ledger", 2014.

Lee Thomas
2016-06-21
Ver 1.1 2016-06-22

Ethereum Network
"Oh look another miner has been successful. Therefore there is little point in continuing with mining this block - we'll start on the next one."
"note - we had to run all the same code on our copies of the EVM to prove this!"

Information required to derive Block Header

Account storage contents Trie
A mapping between integer keys (KEC) and integer values (RLP)

Account, $\sigma[\text{address}]$
RLP data structure
nonce, $\sigma[\text{address}]$
balance, $\sigma[\text{address}]$
storageRoot, $\sigma[\text{address}]$
codeHash, $\sigma[\text{address}]$

World State Trie, σ
A mapping between addresses and account states. Stored as a merkle-patricia tree

Transaction, T
nonce, T_n
gasPrice, T_g
gasLimit, T_l
value, T_v
init, T_i
data, T_d
 $v, r, s, T_n, T_g, T_l, T_v, T_i, T_d$

Transaction Trie, T
A merkle-patricia tree of transactions to include

Transaction Receipts Trie
Index keyed trie

Log Entry, O
Logger's address, O_a
Log topics, O_t
Log data, O_d

Transaction Receipt, $B_t[i]$ (i =transaction no)
post-transaction state, R_s
cumulative gas used, R_g
transaction logs, R_l
bloom filter of log info, R_b

Block, B

Block Header, H or B_H
parentHash, H_p
ommersHash, H_o
beneficiary, H_b
stateRoot, H_s
transactionsRoot, H_t
receiptsRoot, H_r
logsBloom, H_{lb}
difficulty, H_d
number, H_n
gasLimit, H_g
gasUsed, H_{gu}
timestamp, H_{ts}
extraData, H_e
mixHash, H_m
nonce, H_{nc}

Transaction List, B_T

Ommers List, B_U

What are Ethereum Account?

Two types of accounts:

Externally Owned Accounts



0x995a30e66ff7c050c7fdb065e13a83bcf02ba559

- Nonce
- Current Balance
- Code: null
- Storage: null
- Only controlled by private key

Pros:

- Simple
- Light in size
- Predictable

Cons:

- Not customizable
- Cannot see incoming txns

Contract Accounts



- Nonce
- Current Balance
- Contract Code
- Storage
- Controlled only by its code
- Every time it receives a txn, the code is run

Pros:

- Highly Customizable
- Can see all transactions to/from

Cons:

- Can be unsafe
- Bigger blockchain

Transactions vs Messages

Two different ways to interact between accounts:

Transactions:

From EOA to anyone

- **Signature** of the sender
- The recipient of the message (**to**)
- The amount of ether to transfer from the sender to the recipient (**balance**)
- An optional **data** field
- A STARTGAS value
- A GASPRICE value

Messages:

From Contract to other Contracts, think function calls

Contains:

- Same as txn **except** for the signature

They are:

- Never serialized (no transfer over the wire)
- Exists only in the EVM execution environment
- Lead other contract's code to run
- Done using the CALL or DELEGATECALL Opcode

Note: Messages also need gas.

To initiate anything in Ethereum, everything starts from an EOA transaction (which is signed, of course)

Inside an Ethereum Block Header

[illegible]

**** under `transactions`** we only see the transaction hash (for brevity) but really it is a list of transaction objects.

Inside an Ethereum Transaction

[illegible]

Inside an Ethereum Receipt

[illegible]

A little word on **nonces**

Used in many aspects in Ethereum:

- PoW Nonce (random number used with mixHash to prove computation)
- Contract Nonce
- EOA Nonce

Rules for incrementations

- If an externally-owned account sends a transaction, its nonce is incremented **before** execution
- If an externally-owned account creates a contract, its nonce is incremented **before** execution
- If a contract sends a message, **no nonce increments happen**
- If a contract creates a contract, the nonce is incremented **before** the rest of the sub-execution
- The pre-increment nonce is used to determine the contract address
- Nonce increments are never **reverted**

Technical Design

Keccak256 Hashing Algorithm:

The Keccak hashing algorithm is part of the SHA3 family (whereas SHA256, used in bitcoin, is part of the SHA-2 family). It was created by Guido Bertoni, Joan Daemen, Michael Peeters and Gilles Van Assche as part of the NIST hash function competition, organised with the goal to find a new hashing algorithm as it was believed at the time that SHA-2 was potentially reversible.

Ethereum started using it right away and then a new standard, the FIPS-202, changed some constants in the implementation (August 2015) which was called. FIPS-202 was officially called SHA3.

So SHA3 in ethereum **is not** the same SHA3 in the standards!

Ethereum address generation & ECDS choice:

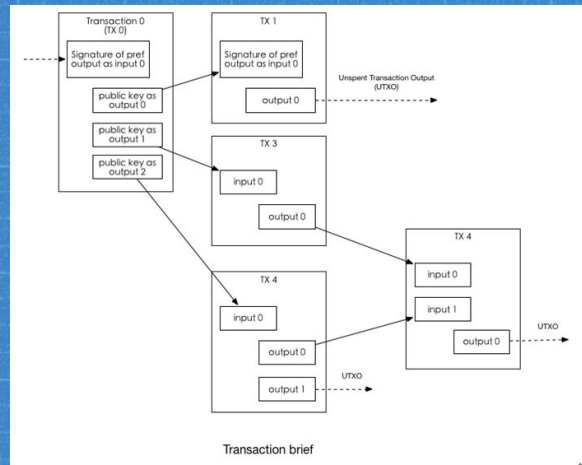
- 1) Same Elliptic curve as bitcoin secp256k1
- 2) Create private key (d)
- 3) Derive public key $Q = d * P$, Q is the public key, P is the base Point (known)
- 4) The address is the **160 rightmost bits** of the Keccak256 hash of the public key
Address = Keccak(publicKey)[96..255]
- 5) Add 0x to the beginning (for Hex purposes) - 42 chars - 21 bytes
- 6) Subsequent addresses (contract created from this account) are the hash of the address with the (pre incremented) nonce.

$$y^2 = x^3 + 7$$

Account Balances, no UTXOs

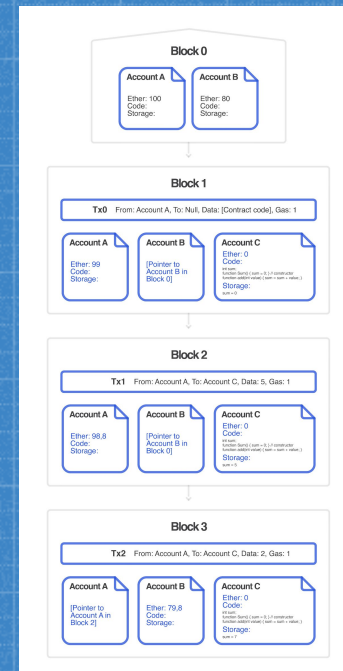
Bitcoin's "State" are all the UTXO's

- Each coin has an owner and a value
- Every referenced input must be valid and not yet spent
- The transaction must have a signature matching the owner of the input for every input
- The total value of the inputs must equal or exceed the total value of the outputs
- The balance is all the UTXO's attached to an owner



Ethereum: Simple Balances

- Space Savings ($1 \text{ UTXO} = 20 + 32 + 8 = 60 \text{ bytes!!}$)
- Greater fungibility
- Simple
- Best for light clients
- Contracts don't need Pub/Priv key pair
- Specific coins are not traceable, only account relationships



The Bitcoin Block Time Problem


Multiple problems are trying to be solved:

- 1) 10 minutes block time is too slow
- 2) BUT we have a way shorter block time, stale rate goes up
- 3) AND stale rate goes up, network **efficiency** and **security** go down
- 4) ALSO there is the problem of centralisation of power through mining pools (>50% computing power in 2014 GHash.io in Bitcoin)

First, how do blocks propagate?

- Blocks are not forwarded directly
- Node sends an **INV** message to other nodes once block has been completely verified
- **INV** message contains:
 - Set of transaction hashes available
 - Set of and block hashes available

Upon noticing a missing element, the receiving node will make a **getData** message containing the missing hashes



Bitcoin Block Propagation time

Block propagation times (Bitcoin):

Stats:

- 6.5 seconds for blocks to propagate to 50% of nodes
- 40 seconds for 95%
- Mean delay of **12.6** seconds

Assuming

- Propagation_time = 12.6 seconds (others in the network are essentially "stagnant" for 12.6s)
- fork_rate = 1.69% (actual)
- Block Time of ~600s

Effective computational power of the network: $1 - 12.6/600 = 0.979$ or 97.9% computing power

Thus, to launch a 51% attack, you only need 48.95% computing power. (w/ block time = 60s → 40% computing power!!)

Note: the block propagation time is dependant on the block size, blocks in ethereum are way smaller

As an aside:

Chances of a txn being confirmed in the real world for Bitcoin

- 63.4% $t < 10$ minutes
- 13.5% $10 < t < 20$ mins
- 0.35% $t > 60$ mins

Other Problem: Mining Pools

Centralization of the network

- Power to censor
- Decision makers in the space

Pool always has an edge

People in the pool will know blocks have been mined faster (since they send messages directly to the participants)

- They thus have less unused computation
- Higher efficiency compared to everyone else (depends on block propagation time && size)

51% Attacks very possible with mining pools

- GHASH.io in 2014

Solution: GHOST Protocol (By Aviv Zohar and Yonatan Sompolinsky)

Big Idea:

- Count the stale blocks (uncles/ommers) as part of the total weight (totalDifficulty in the blockHeader) of the chain

Effects:

- Main chain can hypothetically go down to 5% efficiency
- The attacker will need to outweigh the entire network for a 51% attack (not only mine new blocks but also uncle blocks to get above that difficulty)
- Solves efficiency all the way to 1 second blocks

But it doesn't solve the centralization problem, as a sizeable mining pool could mine 100% of the blocks, (remember, they have the block propagation advantage) thus getting **100% of the rewards**.

Ethereum Solution: **incentivize the uncles!**

If you incentivize the uncles, such that the miners get at maximum their percentage of hashing power as reward, then it doesn't matter so much who is the first one to create a block, what matters is that it is valid. So a 30% miner will get (a bit more than) 30% of the reward.

Logic:

- Every block points to a parent, and can have zero or more uncles.
- Uncle: block w/ valid header which is the child of the parent's parent.
- Actual Block reward: $1 * \text{blockReward} + n * (1/32)$, n being # of uncles included
- Miner of the uncle block gets $7/8 * \text{blockReward}$
- Maximum of 7th generational uncle
- No 2 of the same uncles in the 7 blocks
- Only 2 uncles per block
- "Score" of the block is 0 @ genesis, and then $\text{totalDifficulty} = \text{parentScore} + \text{difficulty} * (1 + \text{uncleCount})$

Gas and fees

Halting Problem! (woot CS!!)

The EVM being Turing complete, we need a way to halt its computation

- When creating a transaction, you provide 2 things:
 - The max amount of computation that you want your program to run (gas)
 - The price per computation you are willing to pay (gasPrice) in Wei (10^{-18} eth)

The total you will be paying for a computation will be $\text{gas} \times \text{gasPrice}$.

Now how does one know how much to pay?

- Every OPCODE has an associated price (which turns out, if not thought of carefully, can lead to DDoS'ing the network)
- 21000 gas for any transaction as a base fee. Covers the cost of the elliptic curve operation to recover the sender's address from the signature as well as the disk bandwidth to store the transaction.
- You can 'fake run' your txn locally to get an estimate at the amount of gas used (not accurate though)

Warning:


- As transactions are atomic, if they run out of gas (Out of Gas exception), everything that happened in the transaction will be reverted (state change etc). Your txn will fail and the miner keeps **ALL** of the gas sent
- If you send too much, you get refunded what wasn't needed

Also, there is a total **block gas limit** that one cannot exceed (currently ~4700000)

- Keeps the blocks small (good for propagation time)
- Keeps execution time fast (remember, miners need to first apply all txns)

Custom Features

- Merkle Patricia Tries
 - Bound to depth of 40
 - $O(\log(n))$ for lookup
 - $O(\log(n))$ for update, insert, delete!!
 - Merkle Proofs for light clients
- RLP serialization
 - Minimalist serialization
 - Easy to implement
 - Safe/guaranteed absolute byte-perfect consistency
- Bloom Filters
 - Efficiently query if txns are members of a block
 - Same for receipt
- EVM
 - Most awesome small computer




The Ethereum Virtual Machine

"The billion dollar computer" -me

EVM

- Stack based virtual machine
- Elements on the stack are 32 byte words
- Same for key/value storage
- Max stack size: 1024
- Call depth limit 1024 (don't do recursion!!)
- Every computational step taken in a program's execution must be paid for up front to prevent DDOS.
- No access to other contract's storage/state
 - Interaction only done through messages where data can be passed (arbitrary length byte array)
- Program execution is sandboxed
- Program execution is fully deterministic and produces identical state transitions for any conforming implementation beginning in an identical state.
- The code for the EVM is in the contract accounts. Upon sending a transaction to a contract, the EVM will do a CALLCODE operation to retrieve the code in the corresponding address

When creating a contract, you actually send a transaction to the "0" address with the relevant data in the data field. This data is not the code of the contract itself, it is the code that will be run by the EVM and **return** the contract code (it will use the SSTORE OPcode to store the bytecode in the "code" field of the account tuple).



Merkle Patricia Tries

Merkle Patricia Tries

Merkle tree

- Tree where all the leaf nodes are hashed together all the way to a root node
- Allows for merkle proofs (authenticate a large dataset using a small amount of data)

Patricia Tries

- PATRICIA: Practical Algorithm To Retrieve Information Coded In Alphanumeric
- Efficient way to store data (great for dictionaries!)
- The key to access the value is actually the path taken from the root node
- Bounded to the maximum length of the key

For txn and receipts trees, they don't need to be MPT's (even though they are):

- Take all the txns in the txn list (they are the **leafs** in the trie)
- Hash all of the txns up to the root node
- Get the root is published on the blockchain

Ethereum's State Implementation, merge both ideas

- Every node parent node is the hash of the children nodes
- The path down to the value is the address (as it is hex, 16 possible values)
- The "value" is the tuple (balance, code, nonce, storage)

Optimizations:

- Different types of nodes:
 - NULL node
 - Leaf node (k,v node)
 - Extension node (k,v node but value is hash of another node)
 - Branch node (list of length 17, first 16 chars are possible hex, final element can hold the k,v node that stops there)
 - Store the node as close to the root as possible

Block Header, H or B_H stateRoot, H_r

Keccak 256-bit hash of the root node of the state trie, after all transactions are executed and finalisations applied

Hash function:

KECCAK256()

Simplified World State, σ

Keys

Values

a	7	1	1	3	5	5	45.0 ETH
a	7	7	d	3	3	7	1.00 WEI
a	7	f	9	3	6	5	1.1 ETH
a	7	7	d	3	9	7	0.12 ETH

World State Trie

ROOT: Extension Node

prefix	shared nibble(s)	next node
0	a7	

Branch Node

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value

Leaf Node

prefix	key-end	value
2	1355	45.0ETH

Extension Node

prefix	shared nibble(s)	next node
0	d3	

Leaf Node

prefix	key-end	value
2	9365	1.1ETH

Prefixes

0 - Extension Node, even number of nibbles
 1□ - Extension Node, odd number of nibbles
 2 - Leaf Node, even number of nibbles
 3□ - Leaf Node, odd number of nibbles
 □ = 1st nibble
 1 nibble = 4 bits

Branch Node


0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value

Leaf Node

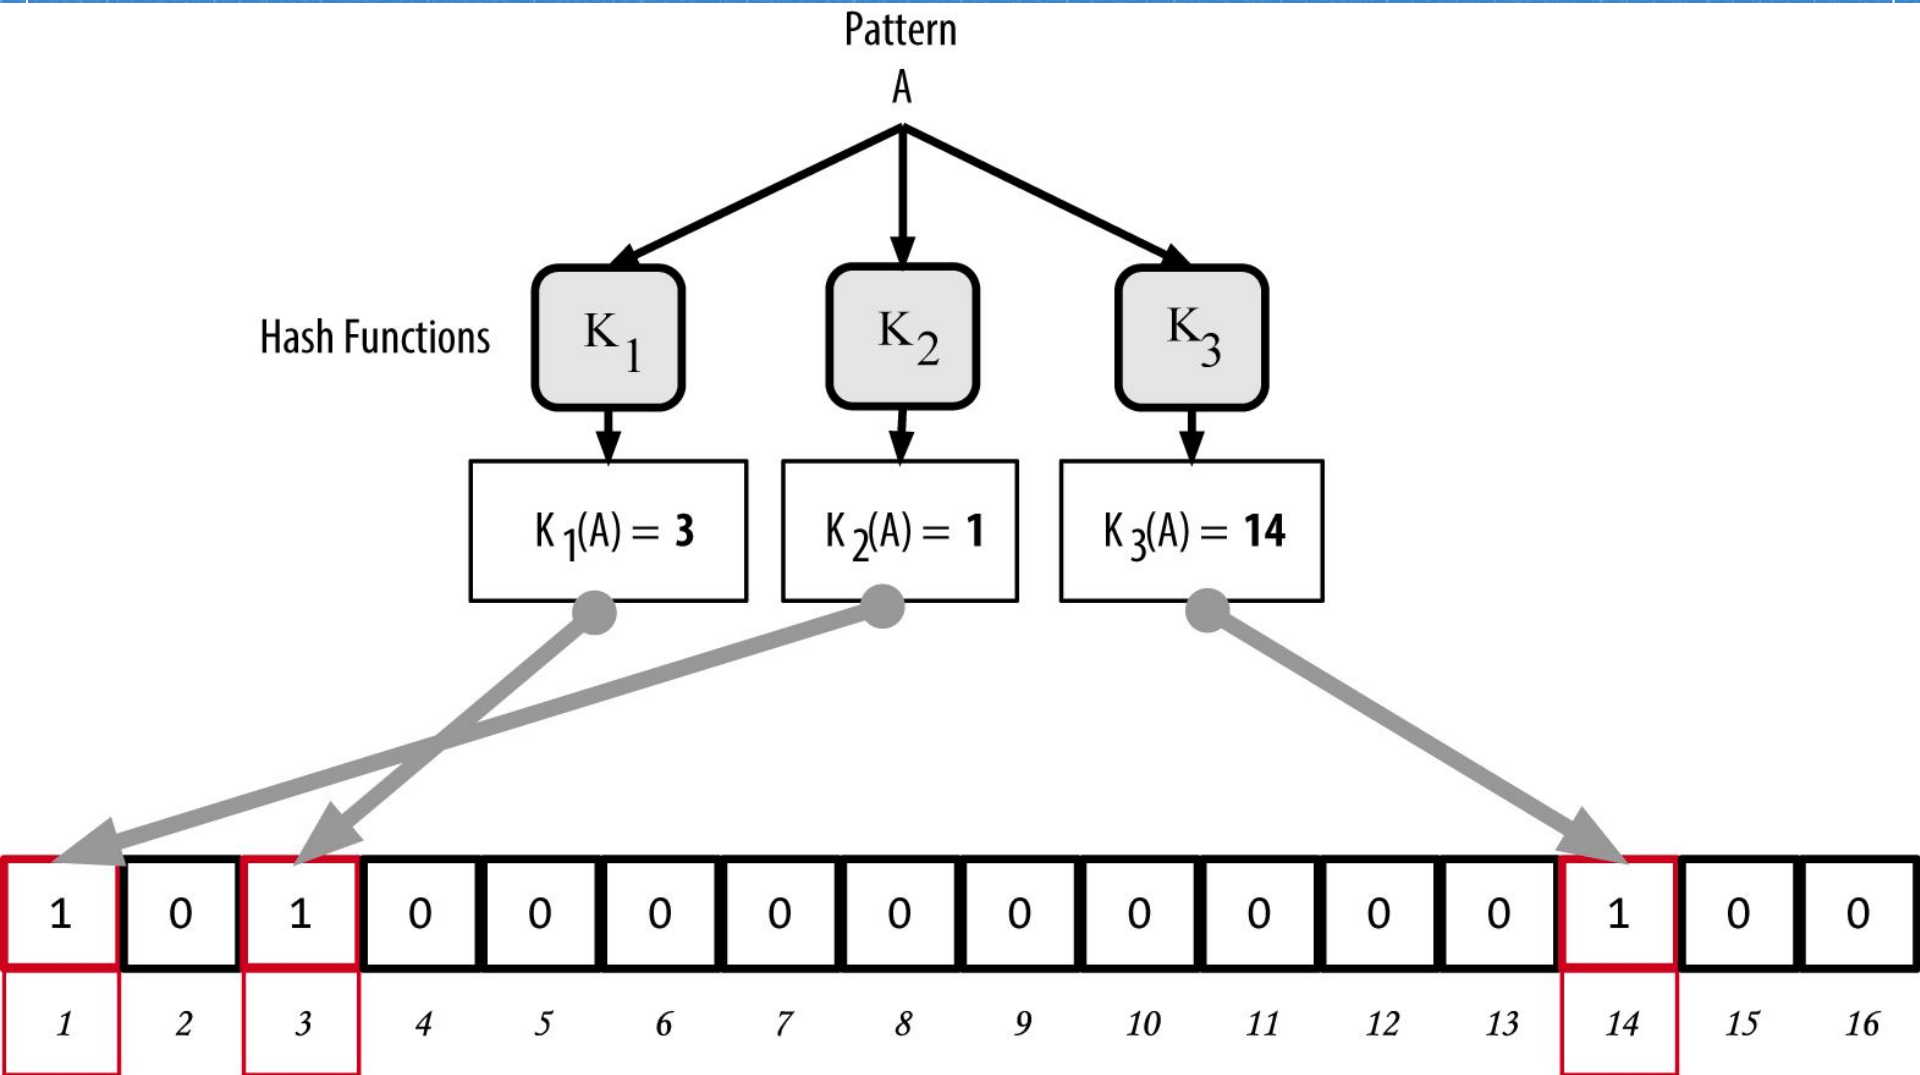
prefix	key-end	value
3□	7	1.00WEI

Leaf Node

prefix	key-end	value
3□	7	0.12ETH



Bloom Filters



Difficulty Bomb

As this difficulty bomb is activated on the network, the mining difficulty will skyrocket and eventually make Ethereum mining unfeasible and extremely unprofitable.

- To switch to PoS
- Supposed to be from block 200 000 (in 2015), now delayed
- Block times could be as high as 14 minutes by 2025.

```
block_diff = parent_diff + parent_diff // 2048 *  
    max(1 - (block_timestamp - parent_timestamp) // 10, -99) +  
    int(2**((block.number // 100000) - 2))
```